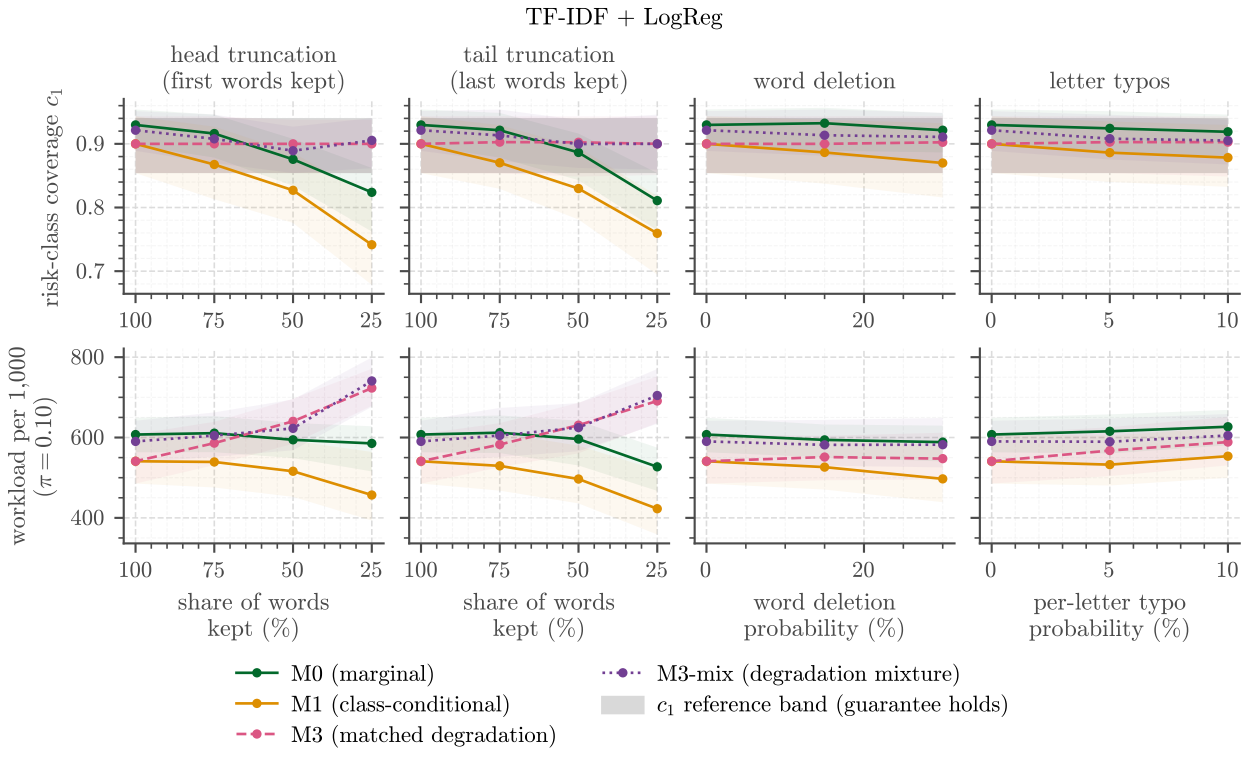
TF-IDF + LogReg
head truncation
(first words kept)
tail truncation
(last words kept)
word deletion
letter typos
risk-class coverage c_1
0.9
0.8
0.7
workload per 1,000
($\pi = 0.10$)
800
600
400
100
75
50
25
0
20
5
10
share of words
kept (%)
word deletion
probability (%)
per-letter typo
probability (%)
M0 (marginal)
M1 (class-conditional)
M3 (matched degradation)
M3-mix (degradation mixture)
c_1 reference band (guarantee holds)